

Million Song dataset

This notebook covers a common supervised learning pipeline, using a subset of the Million Song Dataset (<http://labrosa.ee.columbia.edu/millionsong/>) from the UCI Machine Learning Repository (<https://archive.ics.uci.edu/ml/datasets/YearPredictionMSD>). The goal is to train a linear regression model to predict the release year of a song given a set of audio features.

Part 1: Read and parse the initial dataset

(1a) Load and check the data

The raw data is currently stored in text file. This raw data will be stored as an RDD, with each element of the RDD representing a data point as a comma-delimited string. Each string starts with the label (a year) followed by numerical audio features.

In [33]:

```
# load testing library
from test_helper import Test
import os.path

baseDir = os.path.join('data')
inputPath = os.path.join('cs190', 'millionsong.txt')
fileName = os.path.join(baseDir, inputPath)

numPartitions = 2
rawData = sc.textFile(fileName, numPartitions)
numPoints = rawData.count()
print numPoints
samplePoints = rawData.take(5)
print samplePoints[1]
```

```
6724
2001.0,0.854411946129,0.604124786151,0.593634078776,0.495885413963
,0.266307830936,0.261472105188,0.506387076327,0.464453565511,0.665
798573683,0.542968988766,0.58044428577,0.445219373624
```

(1b) Using LabeledPoint

The labeled training instances are stored using the [LabeledPoint]. The features are the audio attributes and the label is the year of the song.

In [16]:

```
from pyspark.mllib.regression import LabeledPoint
import numpy as np
```

In [43]:

```
def parsePoint(line):
    """Converts a comma separated unicode string into a `LabeledPoint`.

    Args:
        line (unicode): Comma separated unicode string where the first element
            is the label and the
                remaining elements are features.

    Returns:
        LabeledPoint: The line is converted into a `LabeledPoint`, which consists of a label and
            features.
    """
    tokens = line.split(',')
    label, features = tokens[0], tokens[1:]
    return LabeledPoint(label, features)

parsedSamplePoints = [parsePoint(x) for x in samplePoints]
firstPointFeatures = parsedSamplePoints[0].features
firstPointLabel = parsedSamplePoints[0].label
print firstPointFeatures, firstPointLabel

d = len(firstPointFeatures)
print d
```

```
[0.884123733793,0.610454259079,0.600498416968,0.474669212493,0.247
232680947,0.357306088914,0.344136412234,0.339641227335,0.600858840
135,0.425704689024,0.60491501652,0.419193351817] 2001.0
12
```

Visualization 1: Features

First we will load and setup the visualization library. Then we will look at the raw features for 50 data points by generating a heatmap that visualizes each feature on a grey-scale and shows the variation of each feature across the 50 sample data points. The features are all between 0 and 1, with values closer to 1 represented via darker shades of grey.

In [45]:


```
import matplotlib.pyplot as plt
import matplotlib.cm as cm

sampleMorePoints = rawData.take(50)
# You can uncomment the line below to see randomly selected features. These will be randomly
# selected each time you run the cell. Note that you should run this cell with the line commented
# out when answering the lab quiz questions.
# sampleMorePoints = rawData.takeSample(False, 50)

parsedSampleMorePoints = map(parsePoint, sampleMorePoints)
dataValues = map(lambda lp: lp.features.toArray(), parsedSampleMorePoints)

def preparePlot(xticks, yticks, figsize=(10.5, 6), hideLabels=False, gridColor='#999999',
                gridWidth=1.0):
    """Template for generating the plot layout."""
    plt.close()
    fig, ax = plt.subplots(figsize=figsize, facecolor='white', edgecolor='white')
    ax.axes.tick_params(labelcolor='#999999', labelsize='10')
    for axis, ticks in [(ax.get_xaxis(), xticks), (ax.get_yaxis(), yticks)]:
        axis.set_ticks_position('none')
        axis.set_ticks(ticks)
        axis.label.set_color('#999999')
        if hideLabels: axis.set_ticklabels([])
    plt.grid(color=gridColor, linewidth=gridWidth, linestyle='-')
    map(lambda position: ax.spines[position].set_visible(False), ['bottom', 'top', 'left', 'right'])
    return fig, ax

# generate layout and plot
fig, ax = preparePlot(np.arange(.5, 11, 1), np.arange(.5, 49, 1), figsize=(8, 7), hideLabels=True,
                    gridColor='#eeeeee', gridWidth=1.1)
image = plt.imshow(dataValues, interpolation='nearest', aspect='auto', cmap=cm.Greys)
for x, y, s in zip(np.arange(-.125, 12, 1), np.repeat(-.75, 12), [str(x) for x in range(12)]):
    plt.text(x, y, s, color='#999999', size='10')
plt.text(4.7, -3, 'Feature', color='#999999', size='11'), ax.set_ylabel('Observation')
pass
```



(1c) Find the range

The range of the years in the dataset.

In [53]:

```
parsedDataInit = rawData.map(parsePoint)
onlyLabels = parsedDataInit.map(lambda x: x.label).collect()

minYear = min(onlyLabels)
maxYear = max(onlyLabels)
print maxYear, minYear
```

```
2011.0 1922.0
```

(1d) Shift labels

Set the smallest year label 1922 to 0.

In [55]:

```
parsedData = parsedDataInit.map(lambda x: LabeledPoint(x.label-minYear, x.features))

# Should be a LabeledPoint
print type(parsedData.take(1)[0])
# View the first point
print '\n{0}'.format(parsedData.take(1))
```

```
<class 'pyspark.mllib.regression.LabeledPoint'>
```

```
[LabeledPoint(79.0, [0.884123733793,0.610454259079,0.600498416968,
0.474669212493,0.247232680947,0.357306088914,0.344136412234,0.3396
41227335,0.600858840135,0.425704689024,0.60491501652,0.41919335181
7]))]
```


Visualization 2: Shifting labels

The shifted years in the dataset.

In [58]:

```
# get data for plot
newData = (parsedData
            .map(lambda lp: (lp.label, 1))
            .reduceByKey(lambda x, y: x + y)
            .collect())
x, y = zip(*newData)

# generate layout and plot data
fig, ax = preparePlot(np.arange(0, 120, 20), np.arange(0, 120, 20))
plt.scatter(x, y, s=14**2, c='#d6ebf2', edgecolors='#8cbfd0', alpha=0.75)
ax.set_xlabel('Year (shifted)'), ax.set_ylabel('Count')
pass
```



(1e) Training, validation, and test sets

We now split the data into training, validation and test sets using the `[randomSplit method]` with the specified weights and seed to create RDDs storing each of these datasets. Next, cache each of these RDDs.

In [61]:

```
weights = [.8, .1, .1]
seed = 42
parsedTrainData, parsedValData, parsedTestData = parsedData.randomSplit(weights, seed)
parsedTrainData.cache()
parsedValData.cache()
parsedTestData.cache()
nTrain = parsedTrainData.count()
nVal = parsedValData.count()
nTest = parsedTestData.count()

print nTrain, nVal, nTest, nTrain + nVal + nTest
print parsedData.count()
```

```
5371 682 671 6724
6724
```

Part 2: Defining a loss function

(2a) Root mean squared error

We use root mean squared error [RMSE] for evaluation purposes. Implement a function to compute RMSE given an RDD of (label, prediction) tuples, and test out this function on an example.

In [123]:

```
def squaredError(label, prediction):
    """Calculates the the squared error for a single prediction.

Args:
    label (float): The correct value for this observation.
    prediction (float): The predicted value for this observation.

Returns:
    float: The difference between the `label` and `prediction` squared.
    """
    return (float(label)-float(prediction))**2


from operator import add
import math


def calcRMSE(X):
    """Calculates the root mean squared error for an `RDD` of (label, predicti
on) tuples.

Args:
    labelsAndPred (RDD of (float, float)): An `RDD` consisting of (label,
prediction) tuples.

Returns:
    float: The square root of the mean of the squared errors.
    """

    l = X.map(lambda x: float(x[0])).collect()
    p = X.map(lambda x: -float(x[1])).collect()
    c= [sum(x)**2 for x in zip(l, p)]
    return math.sqrt(sum(c)/len(l))


labelsAndPreds = sc.parallelize([(3., 1.), (1., 2.), (2., 2.)])
# RMSE = sqrt[((3-1)^2 + (1-2)^2 + (2-2)^2) / 3] = 1.291
exampleRMSE = calcRMSE(labelsAndPreds)
print exampleRMSE
```

1.29099444874

Part 3: Train (via gradient descent) and evaluate a linear regression model

(3a) Gradient summand

We now apply linear regression, training a model via gradient descent. The gradient descent update for linear regression is:

$$\mathbf{w}_{i+1} = \mathbf{w}_i - \alpha_i \sum_j (\mathbf{w}_i^\top \mathbf{x}_j - y_j) \mathbf{x}_j.$$

where i is the iteration number of the gradient descent algorithm, and j identifies the observation.

First, implement a function that computes the summand for this update, i.e., the summand equals $(\mathbf{w}^\top \mathbf{x} - y)\mathbf{x}$, and test out this function on two examples, using the `DenseVector` method.

In [128]:

```
from pyspark.mllib.linalg import DenseVector
```

In [133]:

```
def gradientSummand(weights, lp):
    """Calculates the gradient summand for a given weight and `LabeledPoint`.

    Note:
        `DenseVector` behaves similarly to a `numpy.ndarray` and they can be u
        sed interchangeably
        within this function. For example, they both implement the `dot` meth
        od.

    Args:
        weights (DenseVector): An array of model weights (betas).
        lp (LabeledPoint): The `LabeledPoint` for a single observation.

    Returns:
        DenseVector: An array of values the same length as `weights`. The gra
        dient summand.
    """
    return (np.dot(weights, lp.features) - lp.label) * lp.features

exampleW = DenseVector([1, 1, 1])
exampleLP = LabeledPoint(2.0, [3, 1, 4])
print np.dot(exampleW, exampleLP.features)
# gradientSummand = (dot([1 1 1], [3 1 4]) - 2) * [3 1 4] = (8 - 2) * [3 1 4]
# = [18 6 24]
summandOne = gradientSummand(exampleW, exampleLP)
print summandOne

exampleW = DenseVector([.24, 1.2, -1.4])
exampleLP = LabeledPoint(3.0, [-1.4, 4.2, 2.1])
summandTwo = gradientSummand(exampleW, exampleLP)
print summandTwo
```

8.0

[18.0, 6.0, 24.0]

[1.7304, -5.1912, -2.5956]

(3b) Use weights to make predictions

Next, implement a `getLabeledPredictions` function that takes in weights and an observation's `LabeledPoint` and returns a (label, prediction) tuple. The prediction is obtained by computing the dot product between weights and an observation's features.

In [135]:

```
def getLabeledPrediction(weights, observation):
    """Calculates predictions and returns a (label, prediction) tuple.

    Note:
        The labels should remain unchanged as we'll use this information to calculate prediction error later.

    Args:
        weights (np.ndarray): An array with one weight for each features in `trainData`.
        observation (LabeledPoint): A `LabeledPoint` that contain the correct label and the features for the data point.

    Returns:
        tuple: A (label, prediction) tuple.
    """
    return (np.dot(weights, observation.features), observation.label)

weights = np.array([1.0, 1.5])
predictionExample = sc.parallelize([LabeledPoint(2, np.array([1.0, .5])),
                                   LabeledPoint(1.5, np.array([.5, .5]))])
labelsAndPredsExample = predictionExample.map(lambda lp: getLabeledPrediction(weights, lp))
print labelsAndPredsExample.collect()
```

```
[(1.75, 2.0), (1.25, 1.5)]
```

(3c) Gradient descent

Next, implement a gradient descent function for linear regression and test out this function on an example.

In [138]:

```
def linregGradientDescent(trainData, numIters):
    """Calculates the weights and error for a linear regression model trained with gradient descent.

    Note:
        `DenseVector` behaves similarly to a `numpy.ndarray` and they can be used interchangeably within this function. For example, they both implement the `dot` method
```

od.

Args:

trainData (RDD of LabeledPoint): The labeled data for use in training the model.

numIters (int): The number of iterations of gradient descent to perform.

Returns:

(np.ndarray, np.ndarray): A tuple of (weights, training errors). Weights will be the final weights (one weight per feature) for the model, and training errors will contain an error (RMSE) for each iteration of the algorithm.

"""

The length of the training data

n = trainData.count()

The number of features in the training data

d = len(trainData.take(1)[0].features)

w = np.zeros(d)

alpha = 1.0

We will compute and store the training error after each iteration

errorTrain = np.zeros(numIters)

for i **in** range(numIters):

Use getLabeledPrediction from (3b) with trainData to obtain an RDD of (label, prediction)

tuples. Note that the weights all equal 0 for the first iteration, so the predictions will

have large errors to start.

labelsAndPredsTrain = trainData.map(**lambda** x: getLabeledPrediction(w,x

))

errorTrain[i] = calcRMSE(labelsAndPredsTrain)

Calculate the `gradient`. Make use of the `gradientSummand` function you wrote in (3a).

Note that `gradient` should be a `DenseVector` of length `d`.

gradient = trainData.map(**lambda** x: gradientSummand(w, x)).reduce(**lambda** x,y: x+y)

Update the weights

alpha_i = alpha / (n * np.sqrt(i+1))

w -= alpha_i*gradient

return w, errorTrain

create a toy dataset with n = 10, d = 3, and then run 5 iterations of gradient descent

note: the resulting model will not be useful; the goal here is to verify that

linregGradientDescent is working properly

exampleN = 10

exampleD = 3

exampleData = (sc

.parallelize(parsedTrainData.take(exampleN))

.map(**lambda** lp: LabeledPoint(lp.label, lp.features[0:exampleD]))

))

print exampleData.take(2)

exampleNumIters = 5

exampleWeights, exampleErrorTrain = linregGradientDescent(exampleData, example

```

NumIters)

print exampleWeights

[LabeledPoint(79.0, [0.884123733793,0.610454259079,0.600498416968]
), LabeledPoint(79.0, [0.854411946129,0.604124786151,0.59363407877
6]))]
[ 48.88110449  36.01144093  30.25350092]

```

(3d) Train the model

Next we train a linear regression model on all of our training data and evaluate its accuracy on the validation set.

In [140]:

```

numIters = 50
weightsLR0, errorTrainLR0 = linregGradientDescent(parsedTrainData,numIters)

labelsAndPreds = parsedValData.map(lambda x: getLabeledPrediction(weightsLR0,x
))
rmseValLR0 = calcRMSE(labelsAndPreds)

print 'Validation RMSE:\n\tBaseline = {0:.3f}\n\tLR0 = {1:.3f}'.format(rmseVal
Base,
                                                                    rmseVal
LR0)

```

```

Validation RMSE:
    Baseline = 21.586
    LR0 = 19.192

```


Visualization 4: Training error

We will look at the log of the training error as a function of iteration. The first scatter plot visualizes the logarithm of the training error for all 50 iterations. The second plot shows the training error itself, focusing on the final 44 iterations.

In [142]:

```
norm = Normalize()
clrs = cmap(np.asarray(norm(np.log(errorTrainLR0))))[:,0:3]


fig, ax = preparePlot(np.arange(0, 60, 10), np.arange(2, 6, 1))
ax.set_ylim(2, 6)
plt.scatter(range(0, numIters), np.log(errorTrainLR0), s=14**2, c=clrs, edgecolors='#888888', alpha=0.75)
ax.set_xlabel('Iteration'), ax.set_ylabel(r'$\log_e(\text{errorTrainLR0})$')
pass
```



In [143]:

```
norm = Normalize()
clrs = cmap(np.asarray(norm(errorTrainLR0[6:]))[:,0:3])

fig, ax = preparePlot(np.arange(0, 60, 10), np.arange(17, 22, 1))
ax.set_ylim(17.8, 21.2)
plt.scatter(range(0, numIters-6), errorTrainLR0[6:], s=14**2, c=clrs, edgecolor='#888888', alpha=0.75)
ax.set_xticklabels(map(str, range(6, 66, 10)))
ax.set_xlabel('Iteration'), ax.set_ylabel(r'Training Error')
pass
```



Part 4: Train using MLlib and perform grid search

(4a) LinearRegressionWithSGD

We can improve the model by adding an intercept, using regularization, and (based on the previous visualization) training for more iterations. First use `LinearRegressionWithSGD` to train a model with L2 regularization and with an intercept. This method returns a `[LinearRegressionModel]`. Next, use the model's `[weights]` and `[intercept]` attributes to print out the model's parameters.

In [144]:

```
from pyspark.mllib.regression import LinearRegressionWithSGD
# Values to use when training the linear regression model
numIters = 500 # iterations
alpha = 1.0 # step
miniBatchFrac = 1.0 # miniBatchFraction
reg = 1e-1 # regParam
regType = 'l2' # regType
useIntercept = True # intercept
```

In [147]:

```
firstModel = LinearRegressionWithSGD.train(parsedTrainData, iterations=numIters,
step=alpha,
miniBatchFraction=1.0, initialWeights=None,
regParam=reg, regType='l2', intercept=True)

# weightsLR1 stores the model weights; interceptLR1 stores the model intercept
weightsLR1 = firstModel.weights
interceptLR1 = firstModel.intercept
print weightsLR1, interceptLR1
```

```
[16.682292427,14.7439059559,-0.0935105608897,6.22080088829,4.01454
261926,-3.30214858535,11.0403027232,2.67190962854,7.18925791279,4.
46093254586,8.14950409475,2.75135810882] 13.3335907631
```

(4b) Predict

Now use the `LinearRegressionModel.predict()`

(<http://spark.apache.org/docs/latest/api/python/pyspark.mllib.html#pyspark.mllib.regression.LinearRegressionModel.predict>) method to make a prediction on a sample point. Pass the features from a `LabeledPoint` into the `predict()` method.

In [162]:

```
samplePoint = parsedTrainData.collect()[0]
print samplePoint.features

samplePrediction = firstModel.predict(samplePoint.features )
print samplePrediction
```

```
[0.884123733793,0.610454259079,0.600498416968,0.474669212493,0.247
232680947,0.357306088914,0.344136412234,0.339641227335,0.600858840
135,0.425704689024,0.60491501652,0.419193351817]
56.8013380112
```

(4c) Evaluate RMSE

Next we evaluate the accuracy of this model on the validation set. Use the `predict()` method to create a `labelsAndPreds` RDD, and then use the `calcRMSE()` function from Part (2a).

In [166]:

```
labelsAndPreds = parsedValData.map(lambda x: (x.label,firstModel.predict(x.features)))  
print labelsAndPreds.top(1)  
rmseValLR1 = calcRMSE(labelsAndPreds)  
  
print ('Validation RMSE:\n\tBaseline = {0:.3f}\n\tLR0 = {1:.3f}' +  
      '\n\tLR1 = {2:.3f}').format(rmseValBase, rmseValLR0, rmseValLR1)
```

```
[(88.0, 61.074692008947231)]
```

```
Validation RMSE:
```

```
    Baseline = 21.586
```

```
    LR0 = 19.192
```

```
    LR1 = 19.691
```

(4d) Grid search

Perform grid search to find a good regularization parameter. We try regParam values 1e-10, 1e-5, and 1.

In [168]:

```
bestRMSE = rmseValLR1
bestRegParam = reg
bestModel = firstModel

numIters = 500
alpha = 1.0
miniBatchFrac = 1.0
for reg in [1e-10, 1e-5, 1]:
    model = LinearRegressionWithSGD.train(parsedTrainData, numIters, alpha,
                                          miniBatchFrac, regParam=reg,
                                          regType='l2', intercept=True)

    labelsAndPreds = parsedValData.map(lambda lp: (lp.label, model.predict(lp.
features)))
    rmseValGrid = calcRMSE(labelsAndPreds)
    print rmseValGrid

    if rmseValGrid < bestRMSE:
        bestRMSE = rmseValGrid
        bestRegParam = reg
        bestModel = model
rmseValLRGrid = bestRMSE

print ('Validation RMSE:\n\tBaseline = {0:.3f}\n\tLR0 = {1:.3f}\n\tLR1 = {2:.3
f}\n' +
      '\tLRGrid = {3:.3f}').format(rmseValBase, rmseValLR0, rmseValLR1, rmseV
allLRGrid)
```

17.0171700716

17.0175981807

23.8007746698

Validation RMSE:

Baseline = 21.586

LR0 = 19.192

LR1 = 19.691

LRGrid = 17.017

Visualization 5: Best model's predictions


Next, we create a visualization 'Predicted vs. actual' fusing the predictions from the best model from Part (4d) on the validation dataset. Specifically, we create a color-coded scatter plot visualizing tuples storing i) the predicted value from this model and ii) true label.

In [170]:

```
predictions = np.asarray(parsedValData
                           .map(lambda lp: bestModel.predict(lp.features))
                           .collect())
actual = np.asarray(parsedValData
                    .map(lambda lp: lp.label)
                    .collect())
error = np.asarray(parsedValData
                   .map(lambda lp: (lp.label, bestModel.predict(lp.features)))
                   .map(lambda (l, p): squaredError(l, p))
                   .collect())

norm = Normalize()
clrs = cmap(np.asarray(norm(error)))[:,0:3]

fig, ax = preparePlot(np.arange(0, 120, 20), np.arange(0, 120, 20))
ax.set_xlim(15, 82), ax.set_ylim(-5, 105)
plt.scatter(predictions, actual, s=14**2, c=clrs, edgecolors='#888888', alpha=
0.75, linewidths=.5)
ax.set_xlabel('Predicted'), ax.set_ylabel(r'Actual')
pass
```



Visualization 6: Hyperparameter heat map

Next, we perform a visualization of hyperparameter search using a larger set of hyperparameters (with precomputed results). Specifically, we create a heat map where the brighter colors correspond to lower RMSE values. The first plot has a large area with brighter colors. In order to differentiate within the bright region, we generate a second plot corresponding to the hyperparameters found within that region.

In [184]:


```
from matplotlib.colors import LinearSegmentedColormap

# Saved parameters and results, to save the time required to run 36 models
numItersParams = [10, 50, 100, 250, 500, 1000]
regParams = [1e-8, 1e-6, 1e-4, 1e-2, 1e-1, 1]
rmseVal = np.array([[ 20.36769649,  20.36770128,  20.36818057,  20.4179535
4,  21.09778437,  301.54258421],
[ 19.04948826,  19.0495,  19.05067418,  19.1651772
6,  19.97967727,  23.80077467],
[ 18.40149024,  18.40150998,  18.40348326,  18.5945749
1,  19.82155716,  23.80077467],
[ 17.5609346,  17.56096749,  17.56425511,  17.8844212
7,  19.71577117,  23.80077467],
[ 17.0171705,  17.01721288,  17.02145207,  17.4451057
4,  19.69124734,  23.80077467],
[ 16.58074813,  16.58079874,  16.58586512,  17.1146690
4,  19.6860931,  23.80077467]])

numRows, numCols = len(numItersParams), len(regParams)
rmseVal = np.array(rmseVal)
rmseVal.shape = (numRows, numCols)

fig, ax = preparePlot(np.arange(0, numCols, 1), np.arange(0, numRows, 1), figs
ize=(8, 7), hideLabels=True,
gridWidth=0.)
ax.set_xticklabels(regParams), ax.set_yticklabels(numItersParams)
ax.set_xlabel('Regularization Parameter'), ax.set_ylabel('Number of Iterations
')

colors = LinearSegmentedColormap.from_list('blue', ['#0022ff', '#000055'], gam
ma=.2)
image = plt.imshow(rmseVal, interpolation='nearest', aspect='auto',
cmap = colors)
```



In [185]:


```
# Zoom into the bottom left
numItersParamsZoom, regParamsZoom = numItersParams[-3:], regParams[:4]
rmseValZoom = rmseVal[-3:, :4]

numRows, numCols = len(numItersParamsZoom), len(regParamsZoom)

fig, ax = preparePlot(np.arange(0, numCols, 1), np.arange(0, numRows, 1), figsize=(8, 7), hideLabels=True,
                      gridWidth=0.)
ax.set_xticklabels(regParamsZoom), ax.set_yticklabels(numItersParamsZoom)
ax.set_xlabel('Regularization Parameter'), ax.set_ylabel('Number of Iterations')

colors = LinearSegmentedColormap.from_list('blue', ['#0022ff', '#000055'], gamma=.2)
image = plt.imshow(rmseValZoom, interpolation='nearest', aspect='auto',
                  cmap = colors)

pass
```



Part 5: Add interactions between features

(5a) Add 2-way interactions

Now, we will add features that capture the two-way interactions between our existing features. A function `twoWayInteractions` takes in a `LabeledPoint` and generates a new `LabeledPoint` that contains the old features and the two-way interactions between them.

In [200]:

```
import itertools

def twoWayInteractions(lp):
    """Creates a new `LabeledPoint` that includes two-way interactions.

    Note:
        For features [x, y] the two-way interactions would be [x^2, x*y, y*x,
        y^2] and these
        would be appended to the original [x, y] feature list.

    Args:
        lp (LabeledPoint): The label and features for this observation.

    Returns:
        LabeledPoint: The new `LabeledPoint` should have the same label as `lp`.
        Its features
        should include the features from `lp` followed by the two-way interaction features.
    """
    lb = lp.label
    a = list(itertools.product(lp.features, repeat=2))
    b = [np.prod(i) for i in a]
    ft = np.hstack((lp.features, b))
    return LabeledPoint(lb, ft)

print twoWayInteractions(LabeledPoint(0.0, [2, 3]))

# Transform the existing train, validation, and test sets to include two-way interactions.
trainDataInteract = parsedTrainData.map(twoWayInteractions)
valDataInteract = parsedValData.map(twoWayInteractions)
testDataInteract = parsedTestData.map(twoWayInteractions)

(0.0, [2.0, 3.0, 4.0, 6.0, 6.0, 9.0])
```

(5b) Build interaction model

Now, we build the new model using the old features and the new interaction features.

In [202]:

```
numIters = 500
alpha = 1.0
miniBatchFrac = 1.0
reg = 1e-10

modelInteract = LinearRegressionWithSGD.train(trainDataInteract, numIters, alpha,
                                              miniBatchFrac, regParam=reg,
                                              regType='l2', intercept=True)

labelsAndPredsInteract = valDataInteract.map(lambda lp: (lp.label, modelInteract.predict(lp.features)))
rmseValInteract = calcRMSE(labelsAndPredsInteract)

print ('Validation RMSE:\n\tBaseline = {0:.3f}\n\tLR0 = {1:.3f}\n\tLR1 = {2:.3f}\n\tLRGrid = ' +
      '{3:.3f}\n\tLRInteract = {4:.3f}').format(rmseValBase, rmseValLR0, rmseValLR1,
                                              rmseValLRGrid, rmseValInteract)
```

```
Validation RMSE:
    Baseline = 21.586
    LR0 = 19.192
    LR1 = 19.691
    LRGrid = 17.017
    LRInteract = 15.690
```

(5c) Evaluate interaction model on test data

We now evaluate the new model on the test dataset. This evaluation provides us with an unbiased estimate for how our model will perform on new data. We also compare with a Baseline model where all the predictions are just the average year.

In [204]:

```
labelsAndPredsTest = testDataInteract.map(lambda lp: (lp.label, modelInteract.predict(lp.features)))
rmseTestInteract = calcRMSE(labelsAndPredsTest)

print ('Test RMSE:\n\tBaseline = {0:.3f}\n\tLRInteract = {1:.3f}'
      .format(rmseTestBase, rmseTestInteract))
```

```
Test RMSE:
    Baseline = 22.137
    LRInteract = 16.327
```

In [206]:

```
predictions = np.asarray(valDataInteract
                          .map(lambda lp: modelInteract.predict(lp.features))
                          .collect())
actual = np.asarray(parsedValData
                   .map(lambda lp: lp.label)
                   .collect())
error = np.asarray(parsedValData
                  .map(lambda lp: (lp.label, bestModel.predict(lp.features)))
                  .map(lambda (l, p): squaredError(l, p))
                  .collect())

norm = Normalize()
clrs = cmap(np.asarray(norm(error)))[:,0:3]

fig, ax = preparePlot(np.arange(0, 120, 20), np.arange(0, 120, 20))
ax.set_xlim(15, 82), ax.set_ylim(-5, 105)
plt.scatter(predictions, actual, s=14**2, c=clrs, edgecolors='#888888', alpha=
0.75, linewidths=.5)
ax.set_xlabel('Predicted'), ax.set_ylabel(r'Actual')
pass
```

